

MSA-ASSIGNMENT-PHASE1 notebook file - Yatai Tian

July 8, 2020

```
[1]: import json
import sys
import time
sys.path.append('/home/nbuser/library/')

import pandas as pd
import requests
import matplotlib.pyplot as plt
import seaborn as sns

df = pd.read_csv('Dataset for Assignment.csv')
```

First, we read in the data set into our work space. Let's check the first few rows and see what we're dealing with.

```
[2]: df.head()
```

	Bedrooms	Bathrooms	Address	Land area \
0	5	3.0	106 Lawrence Crescent Hill Park, Auckland	714
1	5	3.0	8 Corsica Way Karaka, Auckland	564
2	6	4.0	243 Harbourside Drive Karaka, Auckland	626
3	2	1.0	2/30 Hardington Street Onehunga, Auckland	65
4	3	1.0	59 Israel Avenue Clover Park, Auckland	601

	CV	Latitude	Longitude	SA1	0-19 years	20-29 years \
0	960000	-37.012920	174.904069	7009770	48	27
1	1250000	-37.063672	174.922912	7009991	42	18
2	1250000	-37.063580	174.924044	7009991	42	18
3	740000	-36.912996	174.787425	7007871	42	6
4	630000	-36.979037	174.892612	7008902	93	27

	30-39 years	40-49 years	50-59 years	60+ years	Suburbs
0	24	21	24	21	Manurewa
1	12	21	15	30	Karaka
2	12	21	15	30	Karaka
3	21	21	12	15	Onehunga
4	33	30	21	33	Clover Park

This is our function for API calling and we store this into our data frame by adding a column called 'Population'.

```
[3]: def population(lat, lon):
      url = 'https://koordinates.com/services/query/v1/vector.json'
      params = {
          'key': '4a7d61ba2b634a08a297cd2a9f5d582d',
          'layer': '104612',
          'x': lon,
          'y': lat
      }
      response = requests.get(url, params = params)

      pop = response.
      json()['vectorQuery']['layers']['104612']['features'][0]['properties']['C18_CURPop']
      return pop

[4]: df['Population'] = df.apply(lambda x: population(x['Latitude'], x['Longitude']),
      axis = 1)
```

We now get the depreciation index for every SA1 location and merge it with our data frame.

```
[5]: depriv_df = pd.read_excel('otago730395.xlsx')

[6]: merge_df = pd.merge(df, depriv_df[['SA12018_code', 'NZDep2018']], left_on =
      'SA1', right_on = 'SA12018_code')

[7]: merge_df.head()
```

```
[7]: Bedrooms  Bathrooms  Address Land area \
0          5         3.0  106 Lawrence Crescent Hill Park, Auckland    714
1          5         3.0           8 Corsica Way Karaka, Auckland    564
2          6         4.0      243 Harbourside Drive Karaka, Auckland    626
3          2         1.0  2/30 Hardington Street Onehunga, Auckland     65
4          3         1.0      59 Israel Avenue Clover Park, Auckland    601
```



```
CV  Latitude  Longitude  SA1  0-19 years  20-29 years \
0  960000 -37.012920  174.904069  7009770      48         27
1  1250000 -37.063672  174.922912  7009991      42         18
2  1250000 -37.063580  174.924044  7009991      42         18
3   740000 -36.912996  174.787425  7007871      42          6
4   630000 -36.979037  174.892612  7008902      93         27
```



```
30-39 years  40-49 years  50-59 years  60+ years  Suburbs  Population \
0          24          21          24          21  Manurewa    174
1          12          21          15          30   Karaka    129
2          12          21          15          30   Karaka    129
3          21          21          12          15  Onehunga    120
```

4	33	30	21	33	Clover Park	231
---	----	----	----	----	-------------	-----

	SA12018_code	NZDep2018
0	7009770	6.0
1	7009991	1.0
2	7009991	1.0
3	7007871	2.0
4	7008902	9.0

Let us check if there any null values, see the types of each variable and check the data with the describe() function.

```
[8]: merge_df.isnull().values.any()
```

```
[8]: True
```

```
[9]: merge_df.dtypes
```

```
[9]: Bedrooms          int64
Bathrooms            float64
Address              object
Land area            object
CV                   int64
Latitude             float64
Longitude            float64
SA1                  int64
0-19 years           int64
20-29 years          int64
30-39 years          int64
40-49 years          int64
50-59 years          int64
60+ years            int64
Suburbs              object
Population           int64
SA12018_code         int64
NZDep2018            float64
dtype: object
```

```
[10]: merge_df.describe()
```

```
[10]:
```

	Bedrooms	Bathrooms	CV	Latitude	Longitude \
count	1051.000000	1049.000000	1.051000e+03	1051.000000	1051.000000
mean	3.777355	2.073403	1.387521e+06	-36.893715	174.799325
std	1.169412	0.992985	1.182939e+06	0.130100	0.119538
min	1.000000	1.000000	2.700000e+05	-37.265021	174.317078
25%	3.000000	1.000000	7.800000e+05	-36.950565	174.720779
50%	4.000000	2.000000	1.080000e+06	-36.893132	174.798575
75%	4.000000	3.000000	1.600000e+06	-36.855789	174.880944
max	17.000000	8.000000	1.800000e+07	-36.177655	175.492424

	SA1	0-19 years	20-29 years	30-39 years	40-49 years	\
count	1.051000e+03	1051.000000	1051.000000	1051.000000	1051.000000	
mean	7.006319e+06	47.549001	28.963844	27.042816	24.125595	
std	2.591262e+03	24.692205	21.037441	17.975408	10.942770	
min	7.001130e+06	0.000000	0.000000	0.000000	0.000000	
25%	7.004416e+06	33.000000	15.000000	15.000000	18.000000	
50%	7.006325e+06	45.000000	24.000000	24.000000	24.000000	
75%	7.008384e+06	57.000000	36.000000	33.000000	30.000000	
max	7.011028e+06	201.000000	270.000000	177.000000	114.000000	

	50-59 years	60+ years	Population	SA12018_code	NZDep2018
count	1051.000000	1051.000000	1051.000000	1.051000e+03	1051.000000
mean	22.615604	29.360609	179.914367	7.006319e+06	5.063749
std	10.210578	21.805031	71.059280	2.591262e+03	2.913471
min	0.000000	0.000000	3.000000	7.001130e+06	1.000000
25%	15.000000	18.000000	138.000000	7.004416e+06	2.000000
50%	21.000000	27.000000	174.000000	7.006325e+06	5.000000
75%	27.000000	36.000000	210.000000	7.008384e+06	8.000000
max	90.000000	483.000000	789.000000	7.011028e+06	10.000000

Since, land area was an 'object', we wanted to convert this to a float so we
 ↳ remove the letters and characters that are not numbers and convert this to a
 ↳ float.

```
[11]: merge_df['Land area'] = merge_df['Land area'].str.extract('(\d+)').astype(float)
```

```
[12]: merge_df.isnull().sum()
#two missing values in bathrooms and one in suburb
#since our data has over 1000 values, I have decided to drop the three rows with
↳ NULL values
```


```
[12]: Bedrooms      0
      Bathrooms     2
      Address       0
      Land area     0
      CV            0
      Latitude      0
      Longitude     0
      SA1           0
      0-19 years    0
      20-29 years   0
      30-39 years   0
      40-49 years   0
      50-59 years   0
      60+ years     0
      Suburbs       1
      Population    0
      SA12018_code  0
```

```
NZDep2018      0
dtype: int64
```

```
[13]: house_df = merge_df.dropna()
      final_df = house_df.drop('SA12018_code', axis = 1)
```

```
[14]: final_df.isnull().sum()
```

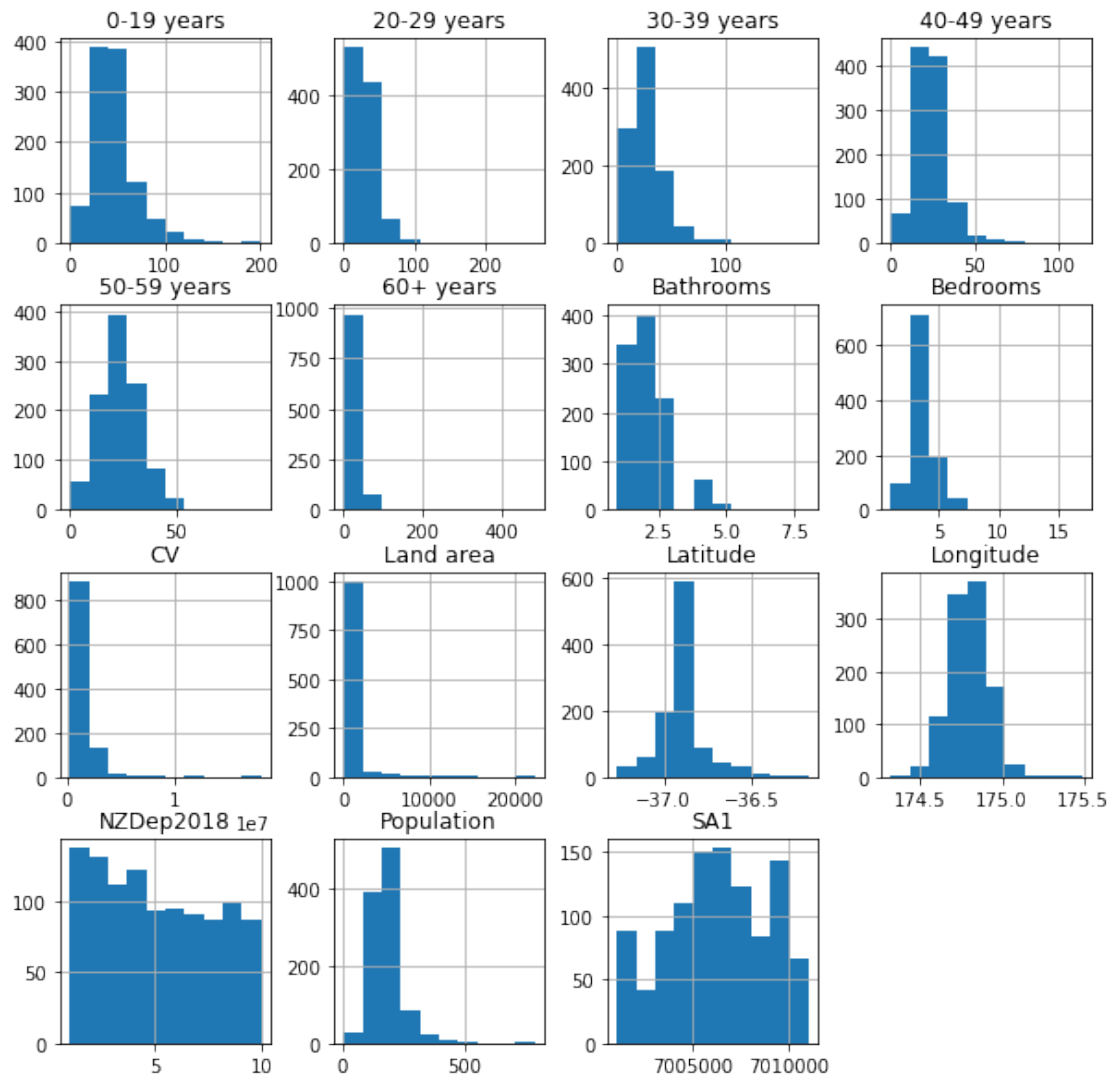
```
[14]: Bedrooms      0
      Bathrooms     0
      Address       0
      Land area     0
      CV            0
      Latitude      0
      Longitude     0
      SA1           0
      0-19 years    0
      20-29 years   0
      30-39 years   0
      40-49 years   0
      50-59 years   0
      60+ years     0
      Suburbs       0
      Population    0
      NZDep2018     0
      dtype: int64
```

Seaborn plots for every numeric variable for an idea on what we're working with. 
↳ Noticeably, CV is heavily left skewed.

```
[15]: final_df.hist(figsize=(10,10))
```

```
[15]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7fcadc71ee48>,
            <matplotlib.axes._subplots.AxesSubplot object at 0x7fcadb2fbf28>,
            <matplotlib.axes._subplots.AxesSubplot object at 0x7fcadb6a4390>,
            <matplotlib.axes._subplots.AxesSubplot object at 0x7fcadbb2a908>],
          [<matplotlib.axes._subplots.AxesSubplot object at 0x7fcadbbd2e80>,
            <matplotlib.axes._subplots.AxesSubplot object at 0x7fcadbed438>,
            <matplotlib.axes._subplots.AxesSubplot object at 0x7fcadb26a9b0>,
            <matplotlib.axes._subplots.AxesSubplot object at 0x7fcadba7af60>],
          [<matplotlib.axes._subplots.AxesSubplot object at 0x7fcadba7af98>,
            <matplotlib.axes._subplots.AxesSubplot object at 0x7fcadb7d8a58>,
            <matplotlib.axes._subplots.AxesSubplot object at 0x7fcadbba2fd0>,
            <matplotlib.axes._subplots.AxesSubplot object at 0x7fcadb730588>],
          [<matplotlib.axes._subplots.AxesSubplot object at 0x7fcadb292b00>,
            <matplotlib.axes._subplots.AxesSubplot object at 0x7fcadbaa80b8>,
            <matplotlib.axes._subplots.AxesSubplot object at 0x7fcadaf9f630>,
            <matplotlib.axes._subplots.AxesSubplot object at 0x7fcadaab4ba8>]],
      dtype=object)
```

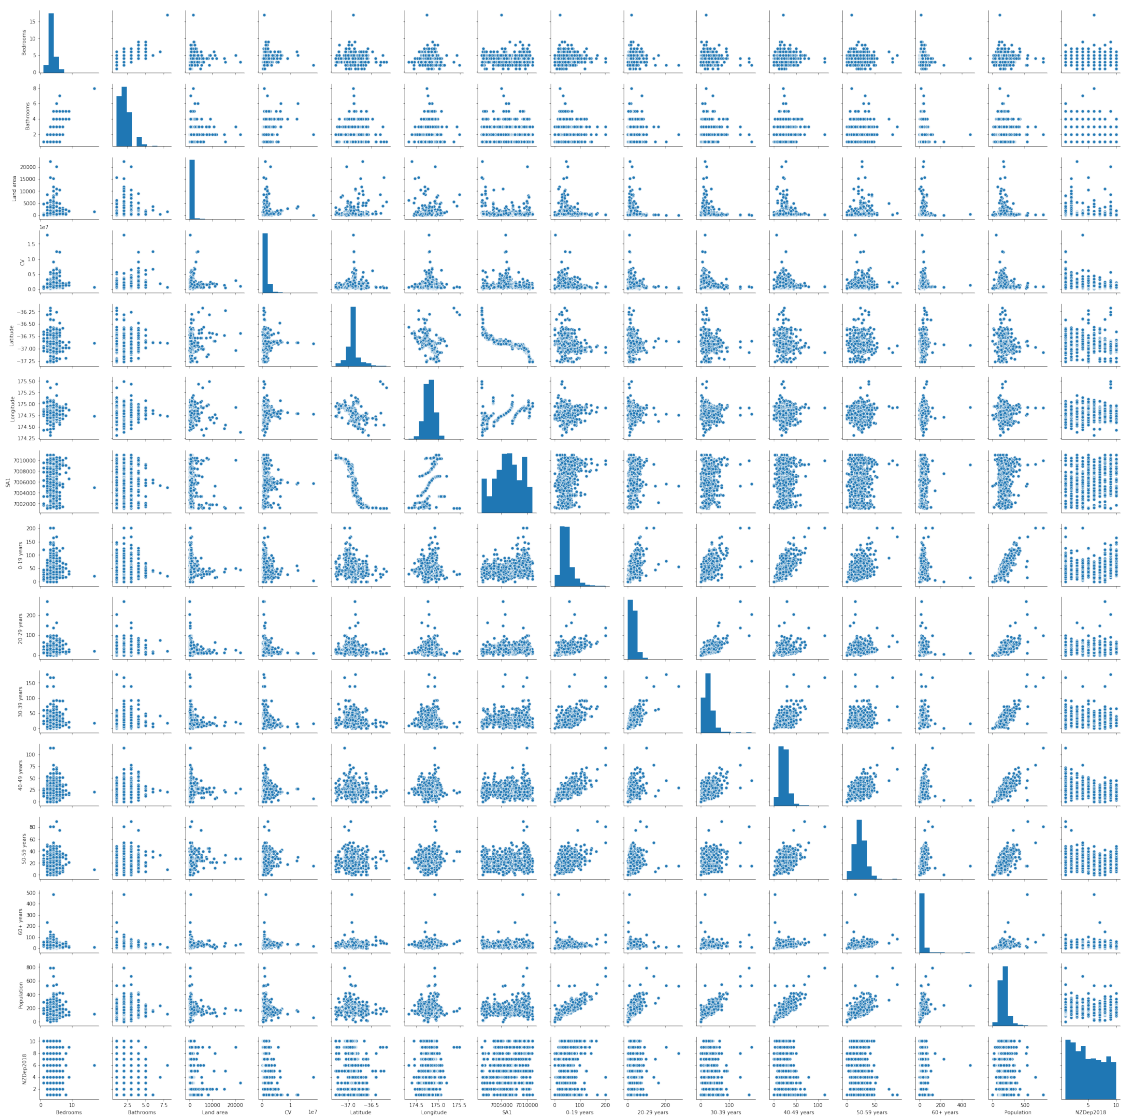
A very mess pairs plot which we will clean up in R.



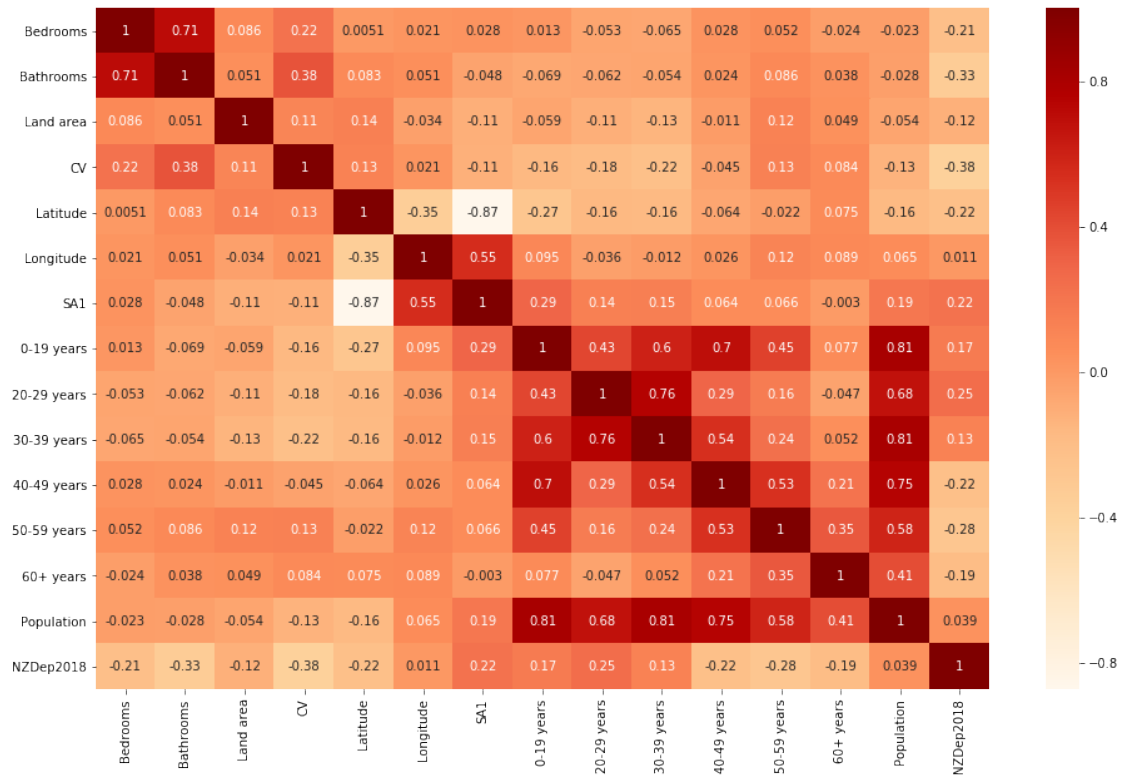
```
[16]: sns.pairplot(final_df, height = 2.0)
```

Finally, a heat map of correlation which will be analysed in depth in R.

```
[16]: <seaborn.axisgrid.PairGrid at 0x7fcadab01198>
```



```
[17]: ax, fig = plt.subplots(figsize=(16,10))
correlation_matrix = final_df.corr()
sns.heatmap(correlation_matrix, annot = True, cmap = 'OrRd')
plt.show()
```



```
[18]: #final_df.to_csv ('final_house.csv', header=True)
```

```
[19]: sns.distplot(final_df['CV'])
#skewed CV so we should log(CV) and use median
```

CV is highly left skewed so we will either $\log(CV)$ or change the distribution in R.
 ↳ further on. I have exported the data set so we can now import it in R.

```
[19]: <matplotlib.axes._subplots.AxesSubplot at 0x7fcab18ff7f0>
```